



JP Hwang

Technical Curriculum Developer
Weaviate

Database Patterns for RAG: Single Collections vs Multi-tenancy



Database Patterns for RAG: Single Collections vs Multi-tenancy:

Prerequisites

- **Docker Desktop**
 - <https://www.docker.com/products/docker-desktop>
- **Python 3.9+**
 - Jupyter Notebook
- **Ollama (Optional) or Cohere API access**
 - If neither is available, the instructor will share a temporary Cohere API key during the session
- **Instructions available online**
 - <https://github.com/weaviate-tutorials/odsc-ai-builders-2025>




Agenda

- 1 RAG overview**
Retrieval augmented generation - the *what* and *why*
- 2 The two patterns**
Introduction to single collection vs multi-tenancy architectures
- 3 Hands-on 1**
Build a RAG system with a single collection type architecture
- 4 Hands-on 2**
Build a RAG system with a multi-tenant architecture
- 5 Recap & comparisons**
When to choose each one?
Implications of choices

RAG: Overview

Large Language Model





Large Language Model

Explain vector
search like I am 5!




prompt



AI outputs

It's like giving each toy a special
"fingerprint" ...

A magic helper that can quickly look
at ALL the toys' "fingerprints" at
once and tell you which ones are
most similar to your query





Generation

Large Language Model

How many world championships
has Max Verstappen won?




prompt



AI outputs

Max Verstappen has won 3 Formula
One World Championships.

He won his first championship in
2021 ... He then dominated the 2022
and 2023 seasons.





Generation

Large Language Model

Who is JP Hwang?



prompt




AI outputs

A South Korean-born, American entrepreneur and businessman.

The founder and CEO of Superb Tickets, an online ticket marketplace.

...






Generation

Large Language Model


Can “hallucinate”



prompt




AI outputs



Generation




Retrieval Augmented Generation







Retrieval Augmented Generation






Retrieval Augmented Generation



Not just RAG...



Agents




Vector Database Patterns

Single Collection




Source Data

Single Collection




Single Collection


- Data schema
- Embedding model definition
- Index configurations
- Replication config
- etc.



Multi-tenant Collection




Multi-tenant Collection





Multi-tenant Collection

- Data schema
- Embedding model definition
- Index configurations
- Replication config
- etc.



Multi-tenant vs Single




- Data schema
- Embedding model definition
- Index configurations
- Replication config
- etc.




Single Collection Pattern

In detail


Single Collection (in detail)




Single Collection (in detail)



Single Collection (in detail)



Single Collection (in detail)





Hands-on section 1!

Single-collection DB patterns



Challenges of scale


Scale: Considerations

- **Object count:** Memory & storage
- **User count:** Data management & compliance
- **Server load:** Distribute load

Managing resource requirements


Scale: Solutions

Growth




Scale: Solutions

Growth




Scale: Solutions

Growth





Scale: Solutions

Growth



Scale: Solutions

Growth



- Single point of failure
- Efficiency
- Upgrades
- Costs



🔍 1024 gib



19 matches

Instance name ▾	On-Demand hourly rate ▲	vCPU ▾	Memory ▾	Storage ▾
x2gd.metal	\$5.344	64	1024 GiB	2 x 1900 SSD
x2gd.16xlarge	\$5.344	64	1024 GiB	2 x 1900 SSD
hpc6id.32xlarge	\$5.70	64	1024 GiB	4 x 3800 NVMe SSD
x2iedn.8xlarge	\$6.669	32	1024 GiB	1 x 950 NVMe SSD
x2idn.16xlarge	\$6.669	64	1024 GiB	1 x 1900 NVMe SSD

AWS Spot pricing, Nov 2024



1024 gib

19 matches

Instance name

On-Demand hourly rate

Storage

x2gd.metal

\$5.344

x 1900 SSD

x2gd.16xlarge

hpc6id.32xlarge

x2iedn.8xlarge

x2idn.16xlarge



$120 * 365 = \$43,800 / \text{year!}$

AWS Spot pricing, Nov 2024



Vector indexing options



Scale: Solutions

Improve efficiency - indexing

Scale: Solutions


Improve efficiency - indexing

- **HNSW** index (default)
- **Flat** index

Scale: Solutions

Indexes - comparison



- **HNSW**: fast + scalable
- **Flat**: tiny footprint; ~100k objs



Scale: Solutions

How to choose?


- **Start with HNSW**
(Tune speed / size / accuracy)
- Multi-tenancy?
 - Try **dynamic**



Scale: Solutions

Improve efficiency - indexing


- **HNSW** index (default)
- **Flat** index
- **Dynamic** index
 - Flat → HNSW @ threshold



Vector quantization

Scale: Solutions


Improve efficiency



Scale: Solutions


Improve efficiency

Compression



Scale: Solutions


Improve efficiency



Compression




Product quantization




Customisable compression

(e.g. 128 floats → 32 bytes: 16x)



Scale: Solutions

Improve efficiency



Compression



Binary quantization

$[-0.1324..., -0.9253...,$
 $0.2389...,$
 $\dots$
 $0.3249..., 0.2390..., -0.4823...]$




$[0, 0, 1, 0, 1, 1,$
 $\dots$
 $0, 1, 0, 1, 1, 0]$

$n \text{ floats} \rightarrow n \text{ bits}$
(32x reduction)

Scale: Solutions

Improve efficiency



Compression



Scalar quantization

$[-0.1324..., -0.9253...,$
 $0.2389...,$
 $\dots$
 $0.3249..., 0.2390..., -0.4823...]$




$[104, 12, 138,$
 $\dots$
 $152, 138, 47]$

$n \text{ floats} \rightarrow n \text{ ints}$
(4x reduction)



Example: Index type & quantization vs memory footprint



HNSW: ~6GB

**HNSW + BQ/PQ:
~2GB**

**Flat (+BQ):
~1GB**



Instance type

vCPU

32

Memory

\$58k / year

vs.

\$23k / year

View



< 1 2 >

Instance name

Demand
hourly rate

vCPU

Memory

Storage

Network
performance

x2iedn.8xlarge

\$6.669

32

1024 GiB

1 x 950 NVMe
SSD

25 Gigabit

x1e.8xlarge

\$6.672

32

976 GiB

1 x 960 SSD

Up to 10 Gigabit

x2gd.8xlarge

\$2.672

32

512 GiB

1 x 1900 SSD

12 Gigabit

AWS Spot pricing, Nov 2024


Scale: Solutions

BQ / PQ / SQ compression

Search **quality** mitigated by over-fetching & rescoring

When to use which?

- Generally, try PQ first
- BQ: model-specific






Multi-Tenancy Pattern


In detail

Single Collection (in detail)




Multi-tenant vs Single

Single-Collection



Multi-Tenant






Hands-on section 2!

Multi-tenant DB patterns


Scale: Solutions

End user growth



Scale: Solutions

End user growth




Challenges faced:

- Performance
- Data isolation
- Compliance

Developed: **Multi-tenancy**

Scale: Solutions

End user growth




Multi-tenancy implementation

- 1000s per node
 - Isolated
 - Active/inactive/offloaded tenants
- tenants

Scale: Solutions

End user growth




Multi-tenancy implementation

- 1000s per node
- **Isolated** (easy deletion & compliance)
- Active/inactive/offloaded tenants

Scale: Solutions



End user growth




Multi-tenancy implementation

- 1000s per node
- Isolated
- **Active/inactive/offloaded tenants** (efficient)


Multi-tenant Collections




Multi-tenant Collections




Multi-tenant Collections




Multi-tenant Collections



Multi-tenant Collections



Multi-tenant Collections



Single collection vs multi-tenant patterns

Single collection patterns

- **Data to be used at the same time:**
 - Search entire dataset
 - HNSW index for scalability
 - Try PQ quantization
- **Example use cases:**
 - E-commerce
 - One Wiki-type knowledge base

Single collection vs multi-tenant patterns

Multi-tenant patterns

- **Isolated datasets for each end user**
 - Subsets searched separately
 - Flat / Dynamic index
 - Try BQ quantization
 - Tenant state management
- **Example use cases:**
 - Online AI-powered apps on private data (wikis, services)



Thank you



weaviate.io



[weaviate/weaviate](https://github.com/weaviate/weaviate)



[JP Hwang](#)